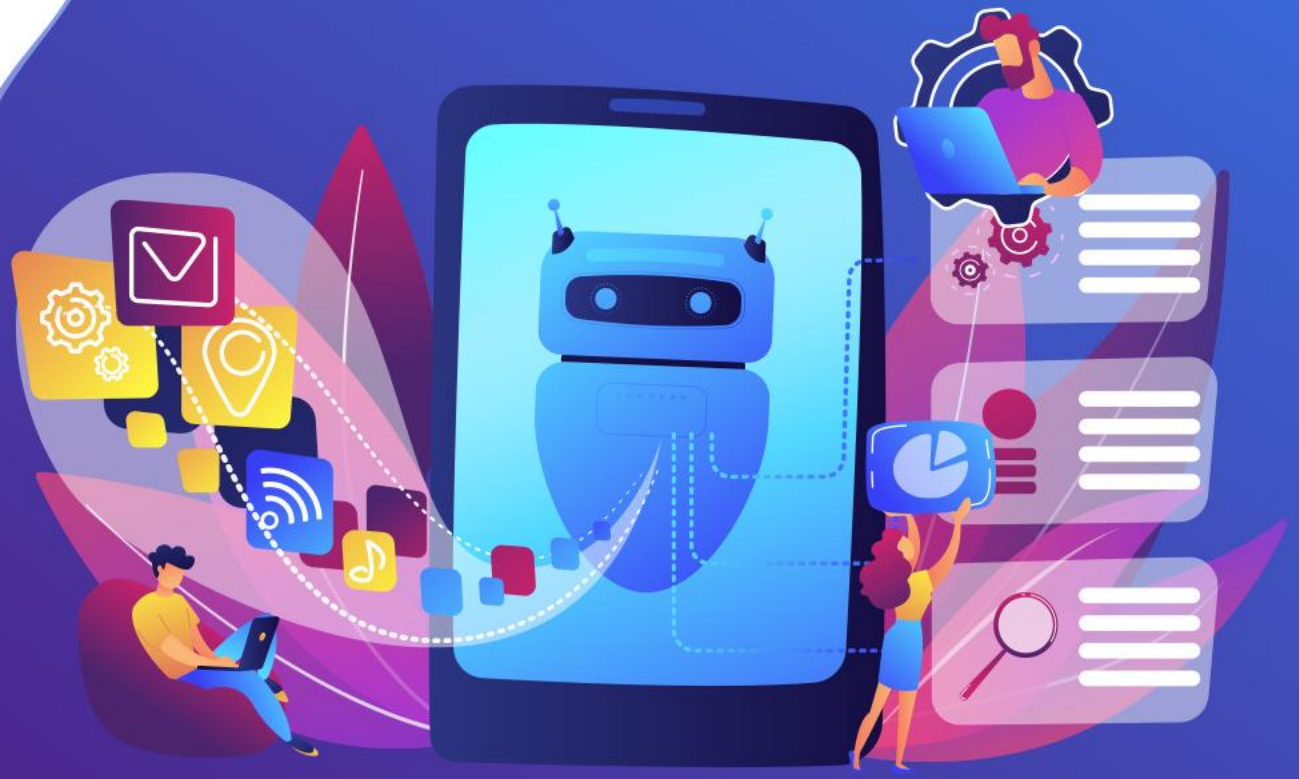

Natural language **processing**

Welcome on board

Lecture 1

Rana Salama

INTRODUCTION TO STATISTICAL NLP

FALL 2020

Outline

- **Course Introduction**
- Course Information
- Deadlines

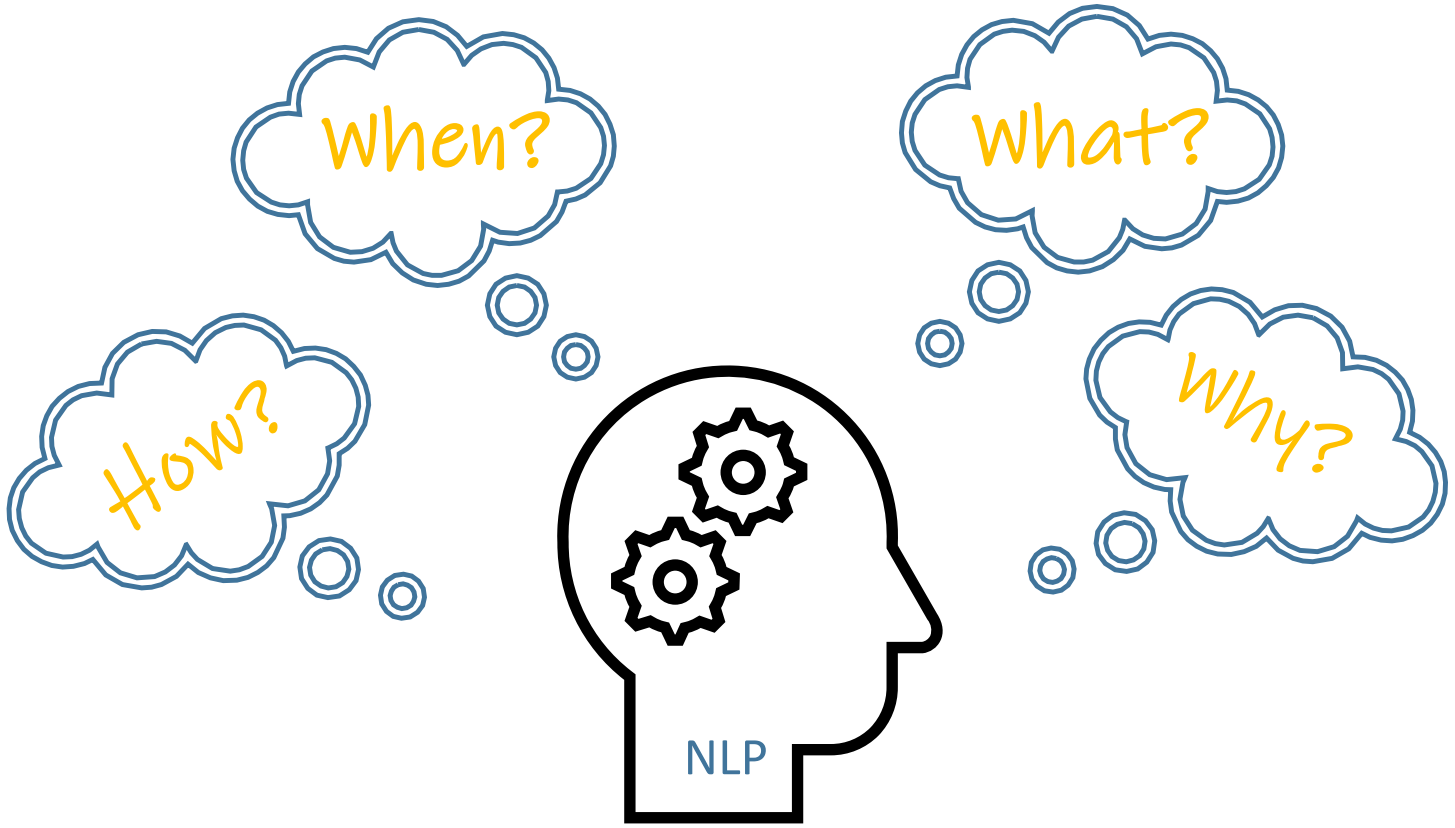

Course Introduction

What is NLP?

- Natural language processing (NLP) is the interdisciplinary field of computer science and linguistics

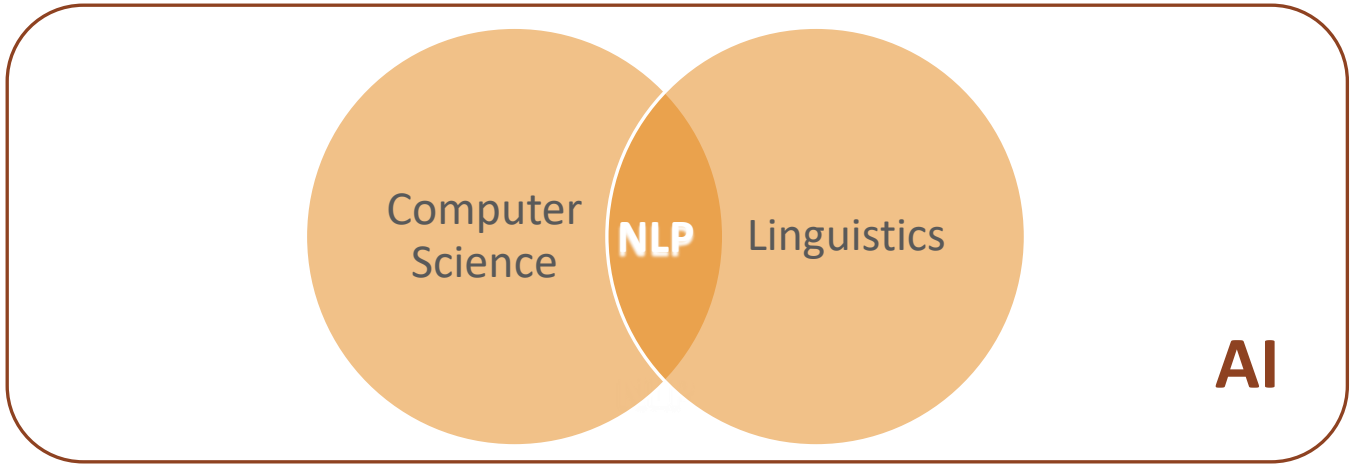

- **Goal:** Have computers *understand* natural language in order to perform useful tasks (human-machine communication or improving human-human communication)

What is NLP?

Natural Language

- Natural language refers to the language used by humans to communicate with each other.

-Human language is diverse and complex. 6500 Languages exist

NLP

Processing

-Computers process natural language input (unstructured data) and produce useful information.

-Programs that can recognize, analyze , and generate text

-NLP refers to processing unstructured data – text in free form (unstructured text).

Natural Language

Language = Words (Dictionary) and Rules (Grammar)

Dictionary: set of words defined in the language; open (dynamic)

Grammar: set of rules which describe what is allowable in a language

-**Classical Grammars**: meant for humans; mainly supported by examples; no (or almost no) formal description tools; cannot be programmed

A complete sentence must include a noun and a verb

"the bird flew"

- **Explicit Grammar**: (CFG, Dependency Grammars, Link Grammars,...) formal description; can be programmed & tested on data (texts)

S-> NP VP

"I prefer a morning flight"

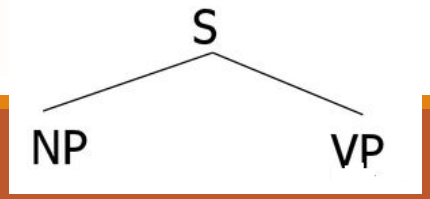

Layers of linguistic analysis

- Linguistic is the science of language.
- Its study include 6 basic levels (more or less explicitly present in most theories):
  - ❑ Phonetics and Phonology: sounds
  - ❑ Morphology : word formation
  - ❑ Syntax : structural relationships between words, sentence formation
  - ❑ Semantics : knowledge of meaning
  - ❑ Pragmatic: connected sentences
- Each level has an input and output representation
- Output from one level is the input to the next (upper) level
- Sometimes levels might be skipped (merged) or split

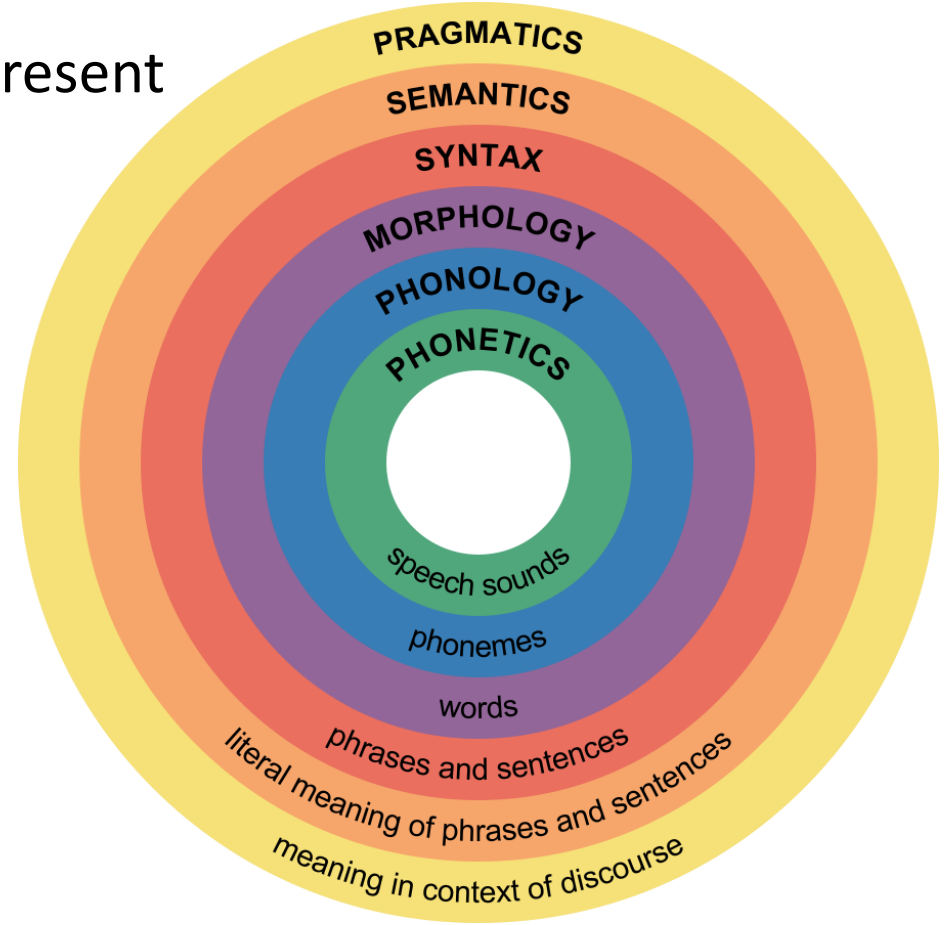

Why NLP?

- An enormous amount of knowledge is now available in machine readable form as unstructured natural language text from different resources (heterogeneous data)
- Going from the largely unstructured languages of the web to useful information
- Conversational agents are becoming an important form of human-computer communication
- Much of human-human communication is now mediated by computers.
- Very cool stuff! And with lots of commercial interest.

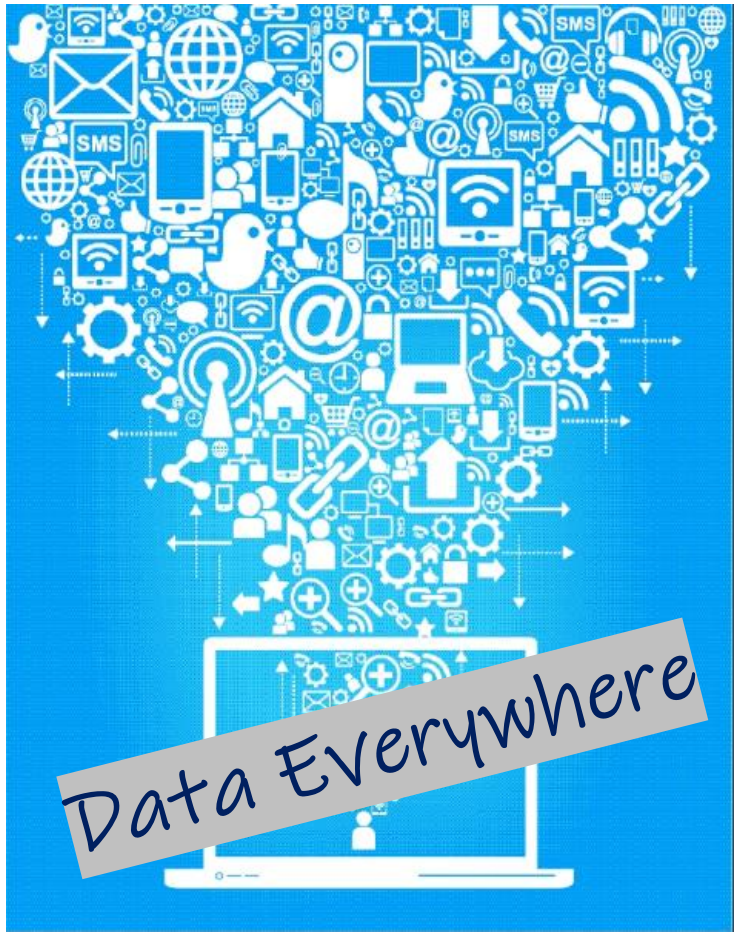

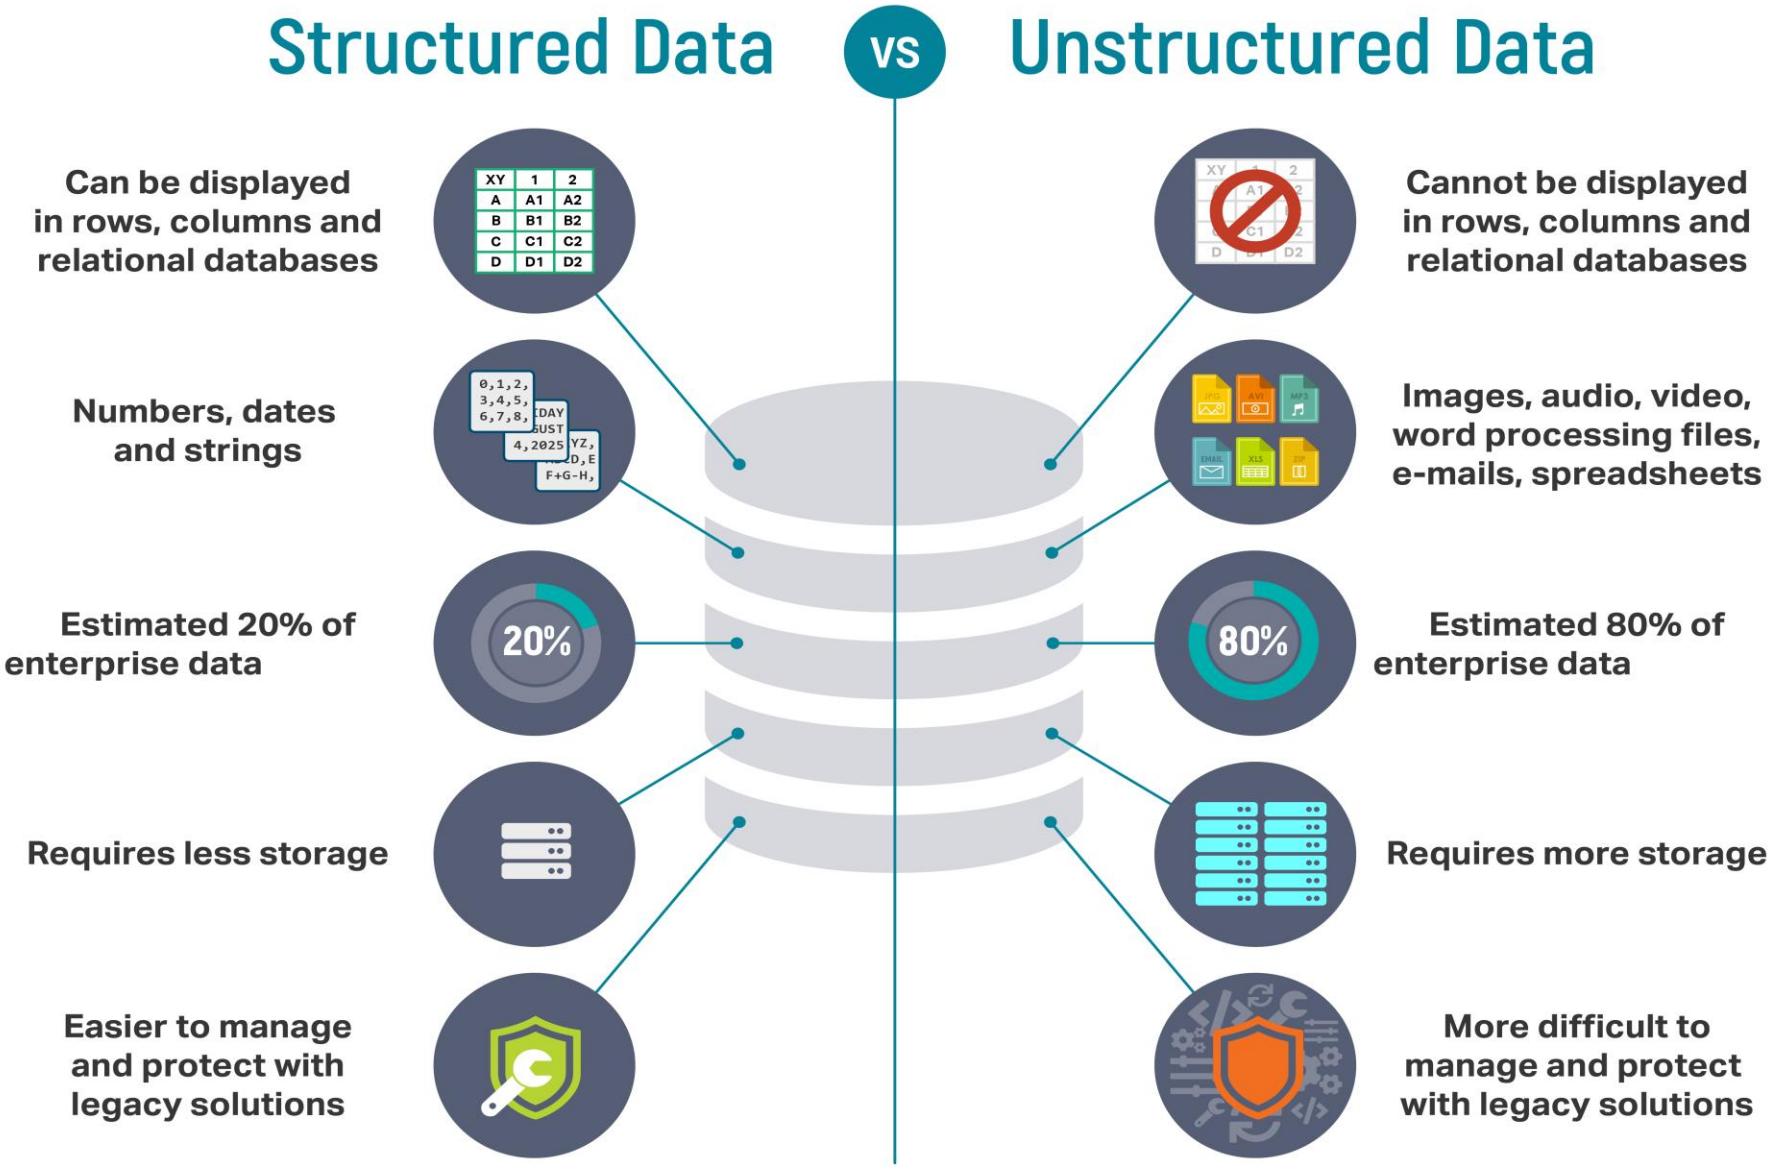

https://www.igneous.io/blog/structured-data-vs-unstructured-data

Why NLP ?

Aoccdrnig to a rscheearch at Cmabrigde Uinervtisy, it deosn't mttaer in waht oredr the ltteers in a wrod are, the olny iprmoatnt tihng is taht the frist and lsat ltteers be at the rghit pclae. The rset can be a toatl mses and you can sitll raed it wouthit porbelm.

Why NLP ?

- People have no trouble understanding language
  - Commonsense knowledge
  - Reasoning capacity
  - Experience
- However, Computers have
  - No commonsense knowledge
  - No reasoning capacity

Unless we teach them!

Why NLP ?

We need computers to :

- Classify text into categories
- Index and search large texts
- Automatic machine translation
- Speech understanding – Understand phone conversations
- Information extraction – Extract useful information from resumes
- Automatic summarization – Condense 1 book into 1 page •
- Question answering
- Knowledge acquisition
- Text generation / dialogs

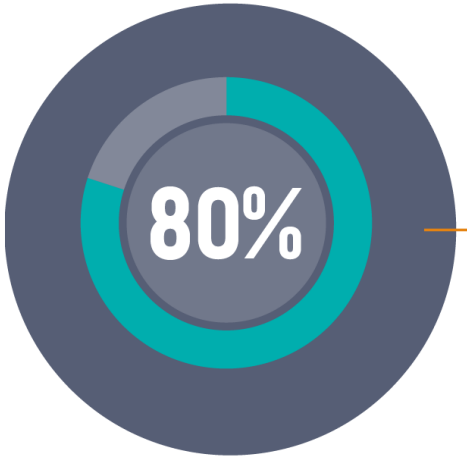

Unstructured Data

NLP

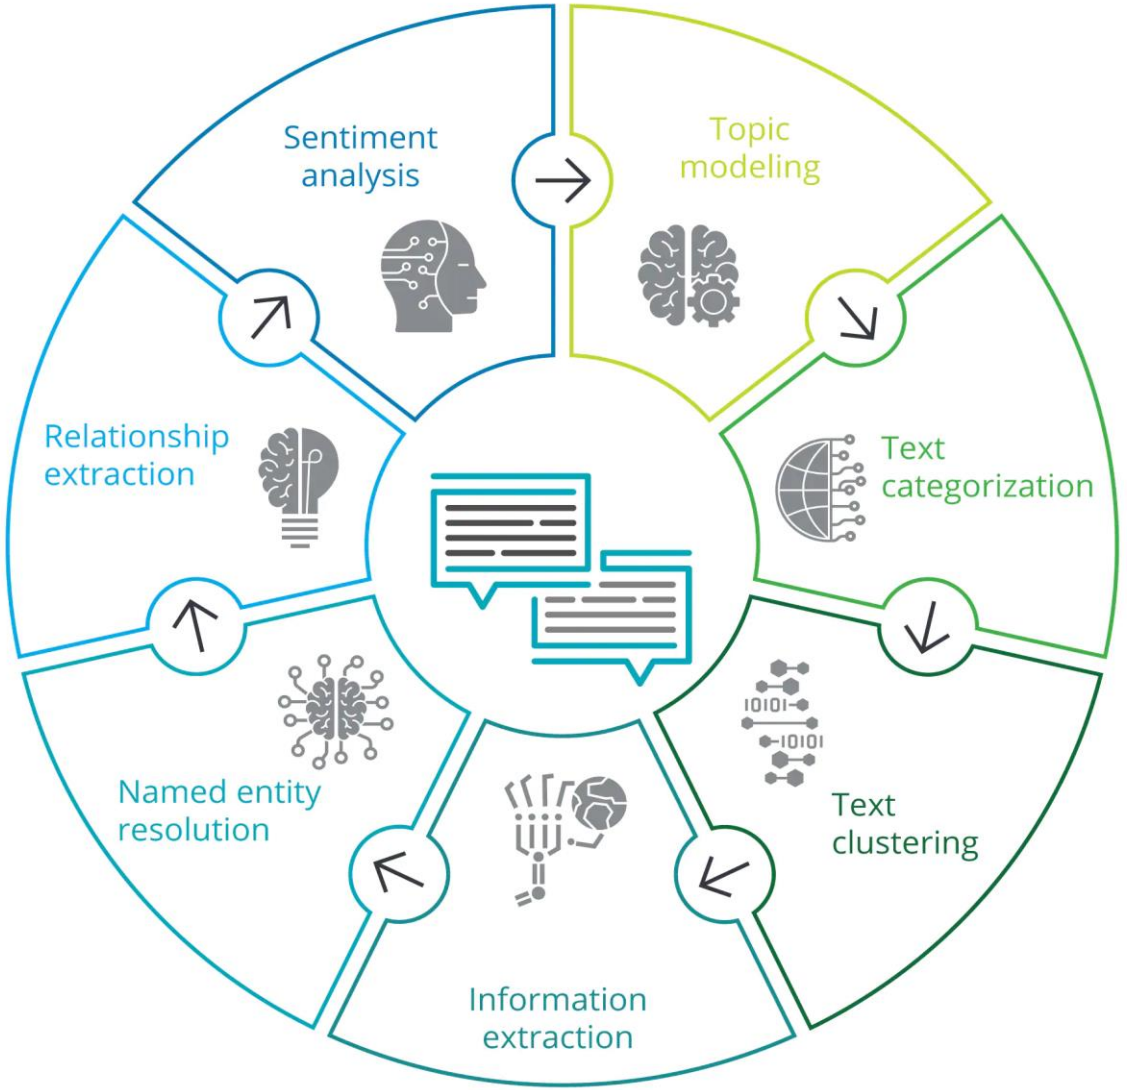

Source: Deloitte analysis.

NLP Applications

NLP Everywhere

USER PERSPECTIVE

Auto-complete

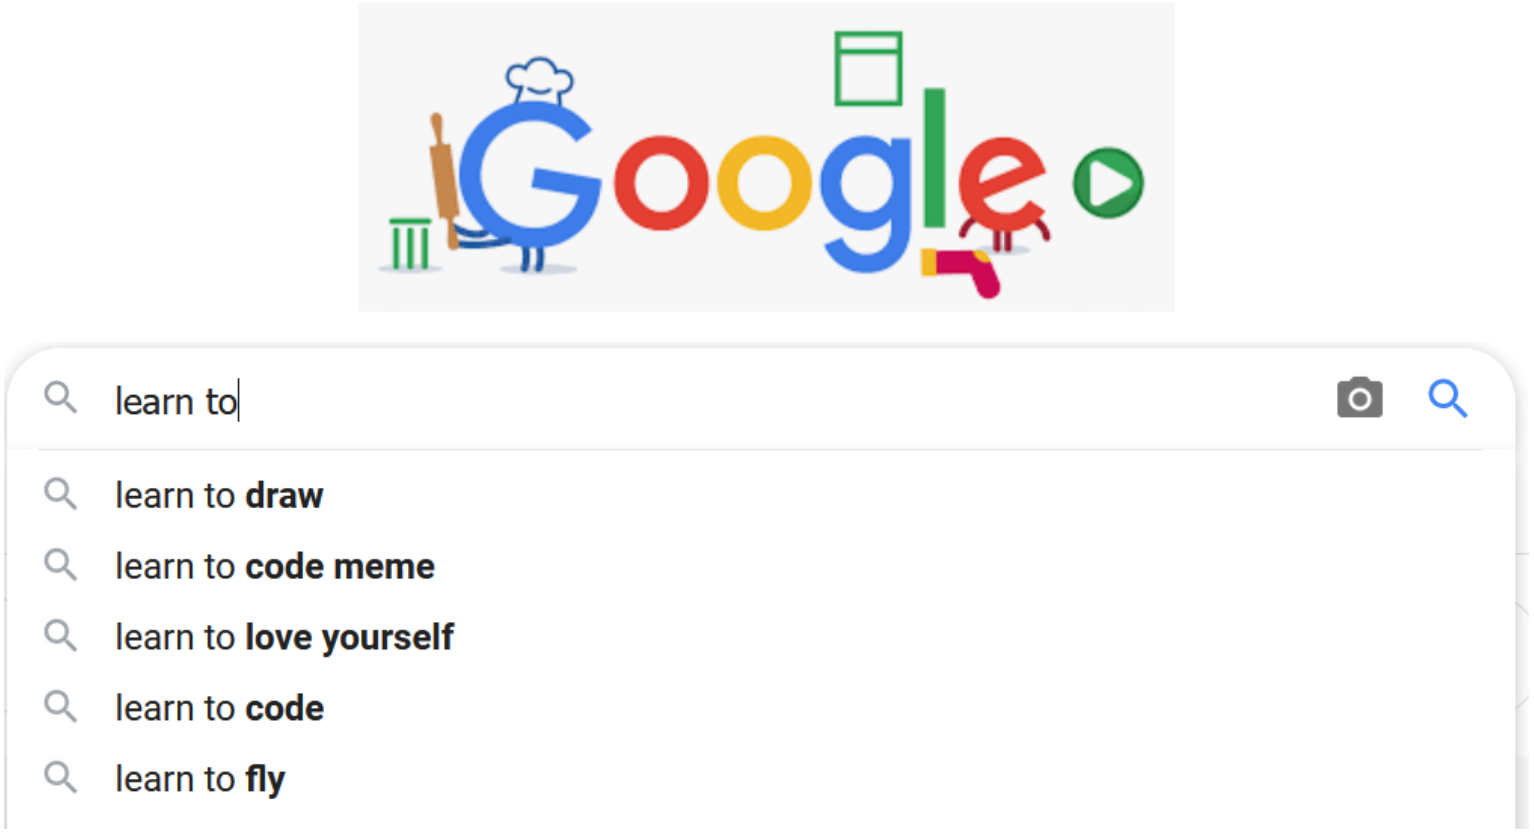

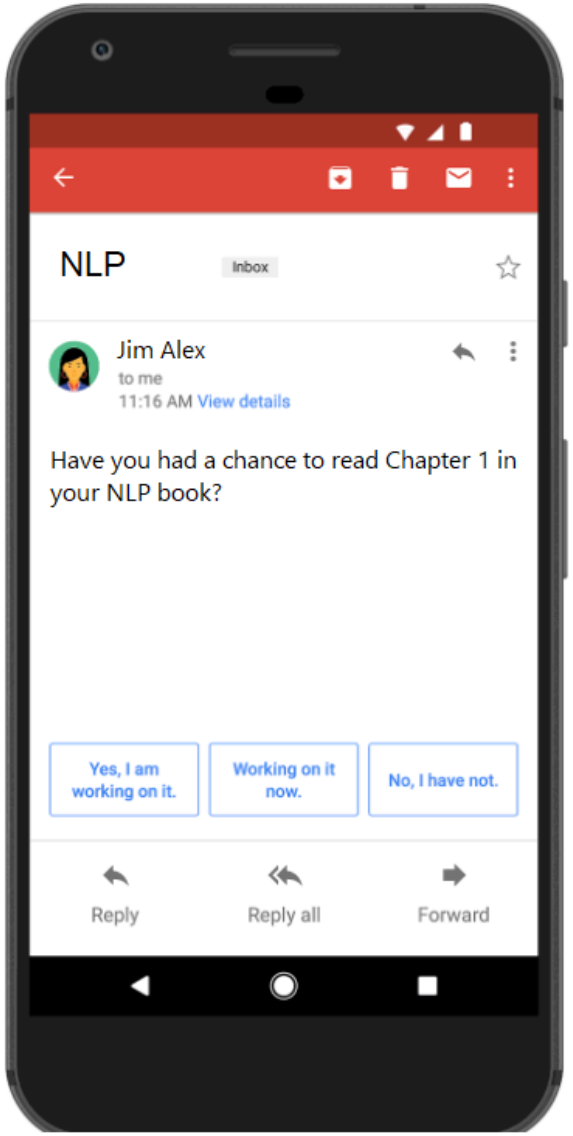

Automatic Replies

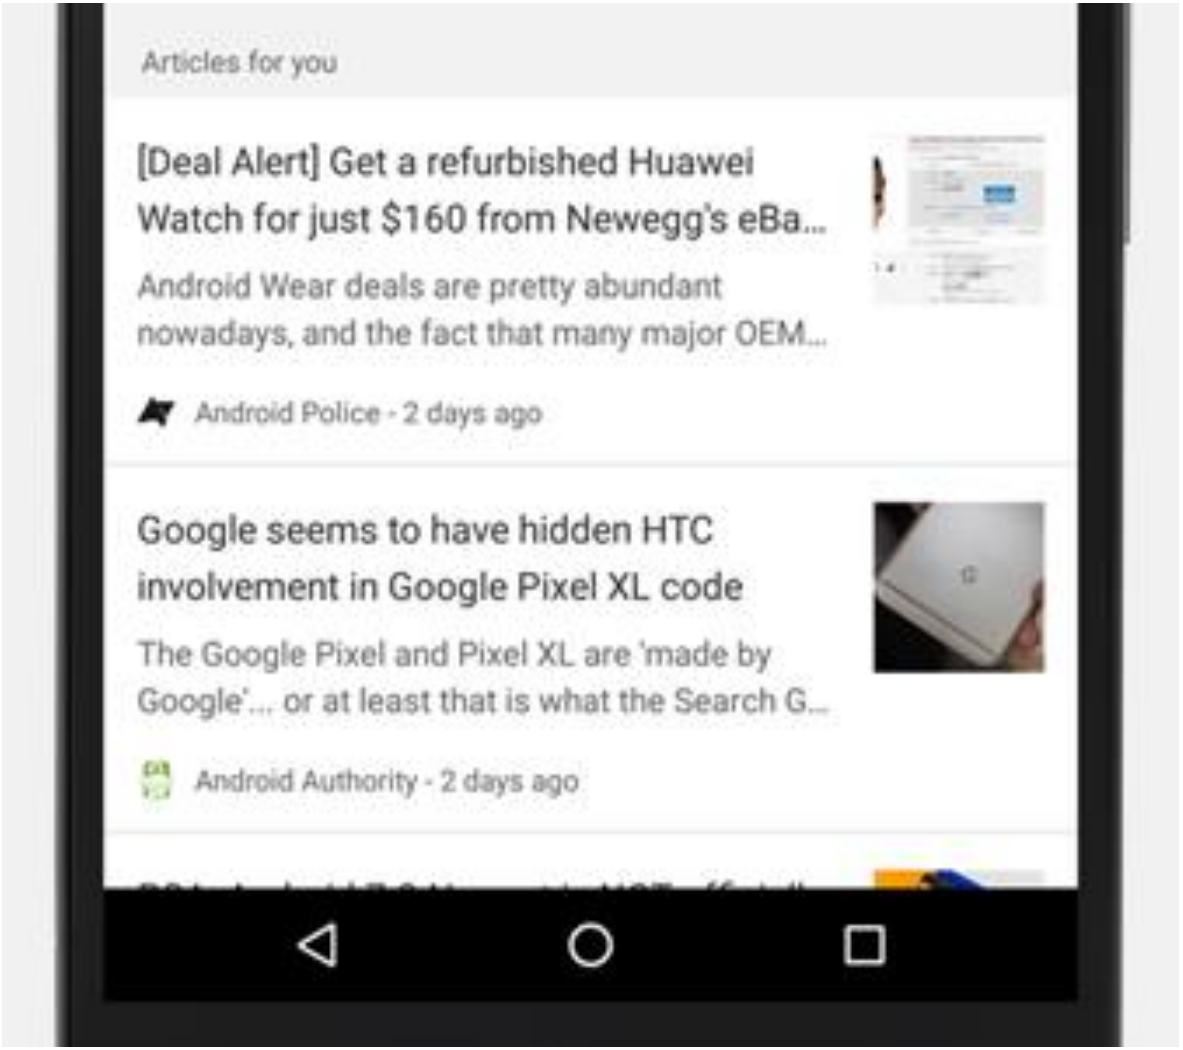

News site's suggested articles

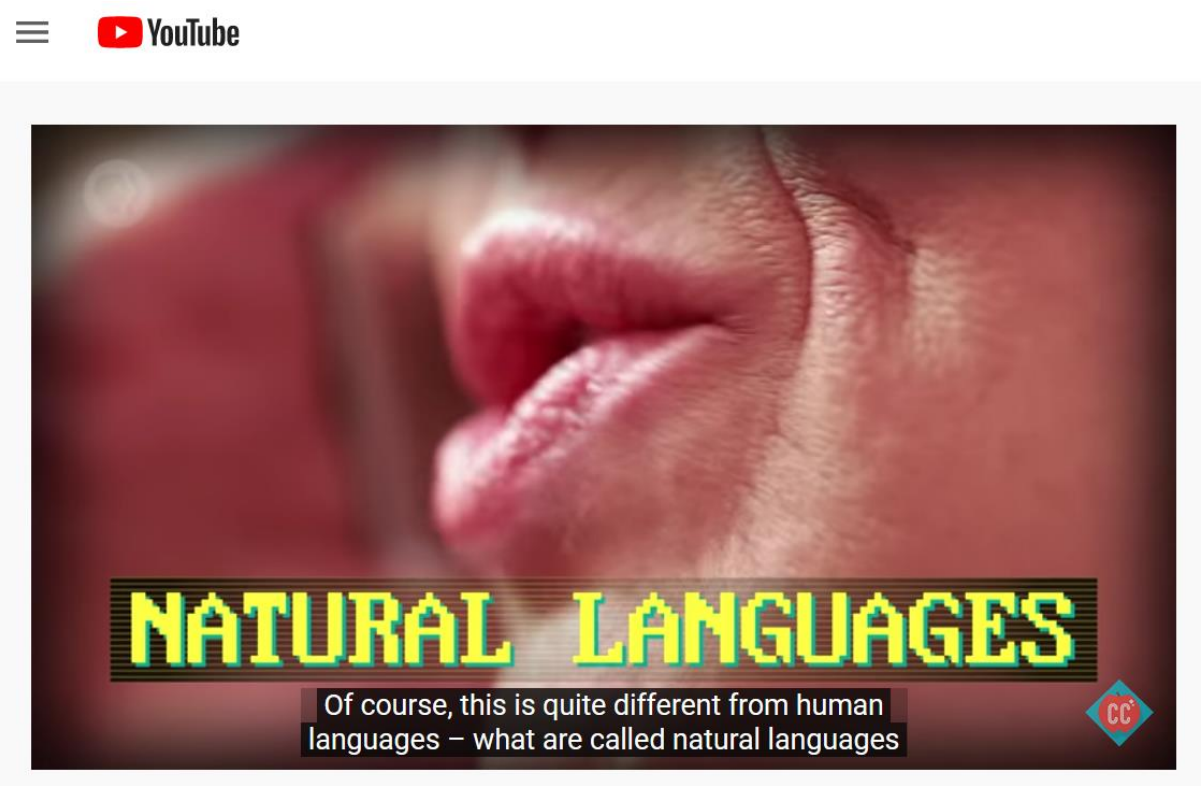

Auto-generated video captions

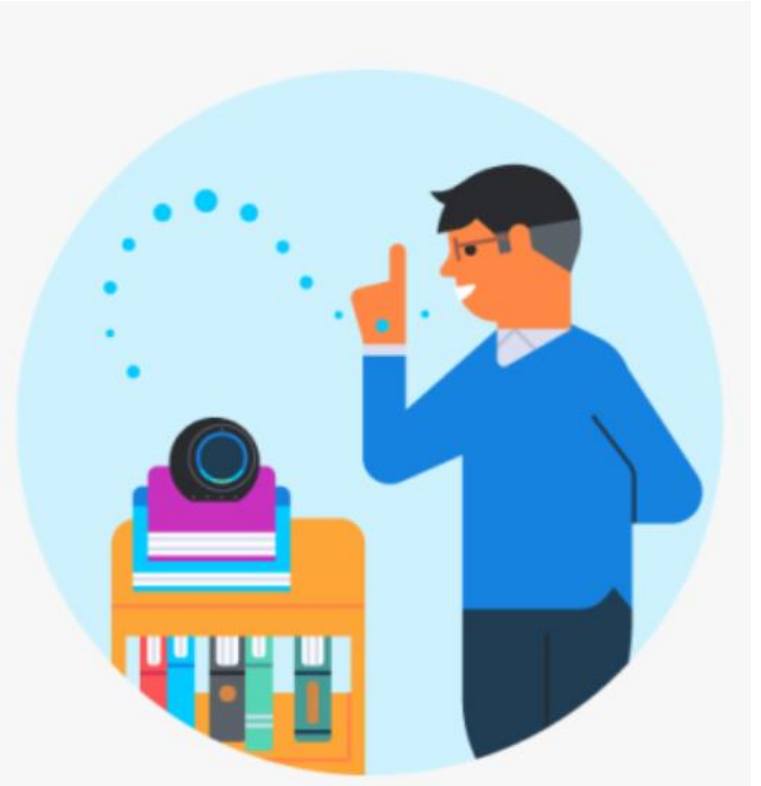

Hey Alexa,..

And more to come

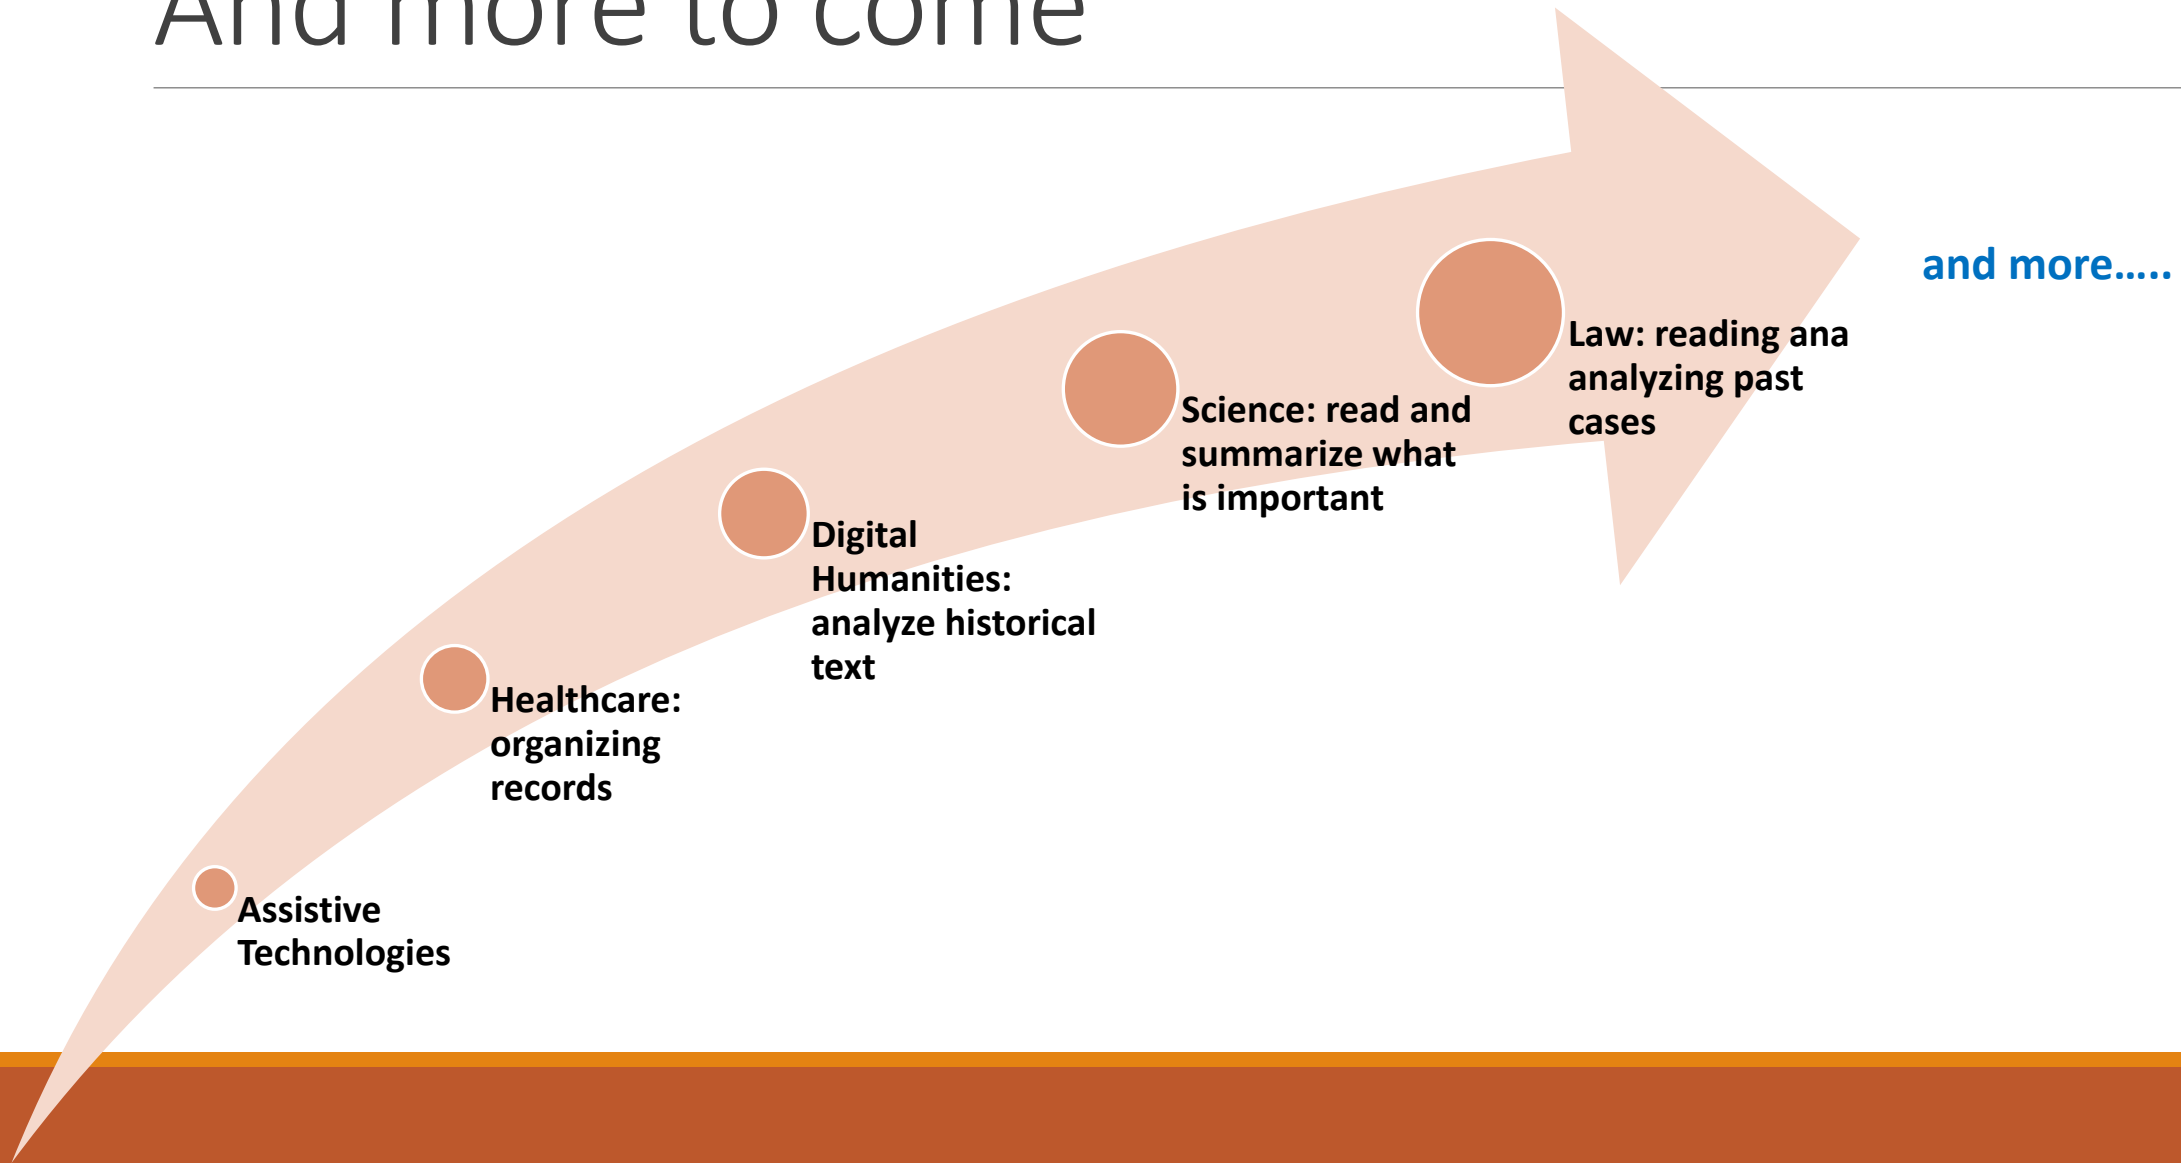

NLP Applications

A CLOSER LOOK

Machine Translation

- **Goal**: Computers understand and translate between one language and another: e.g. Google Translate.

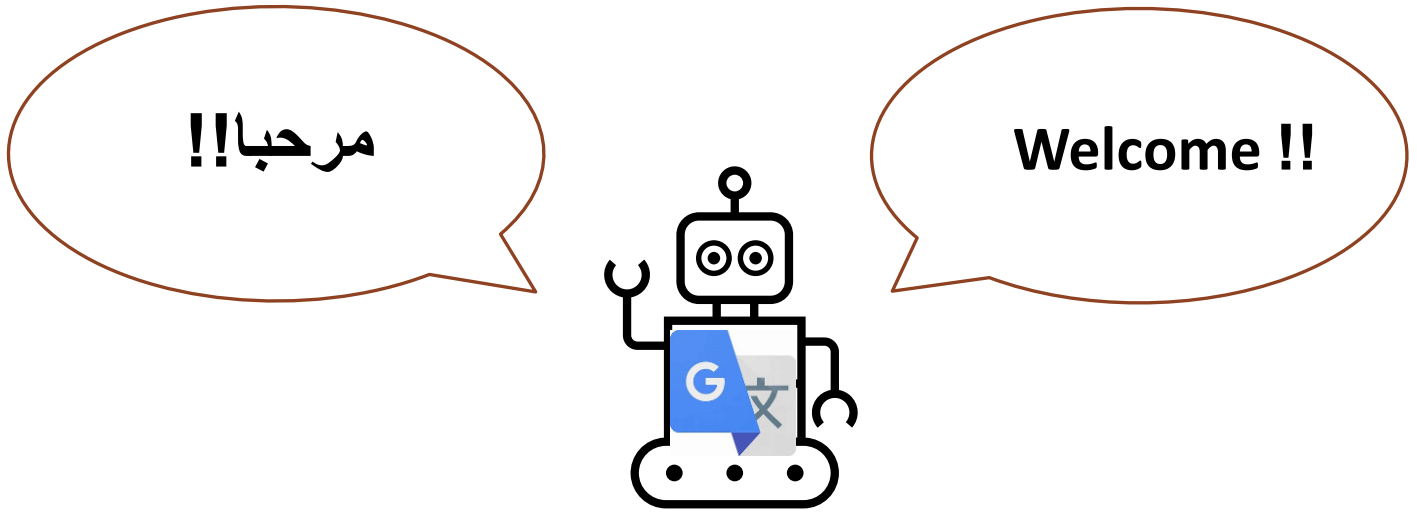

Machine Translation

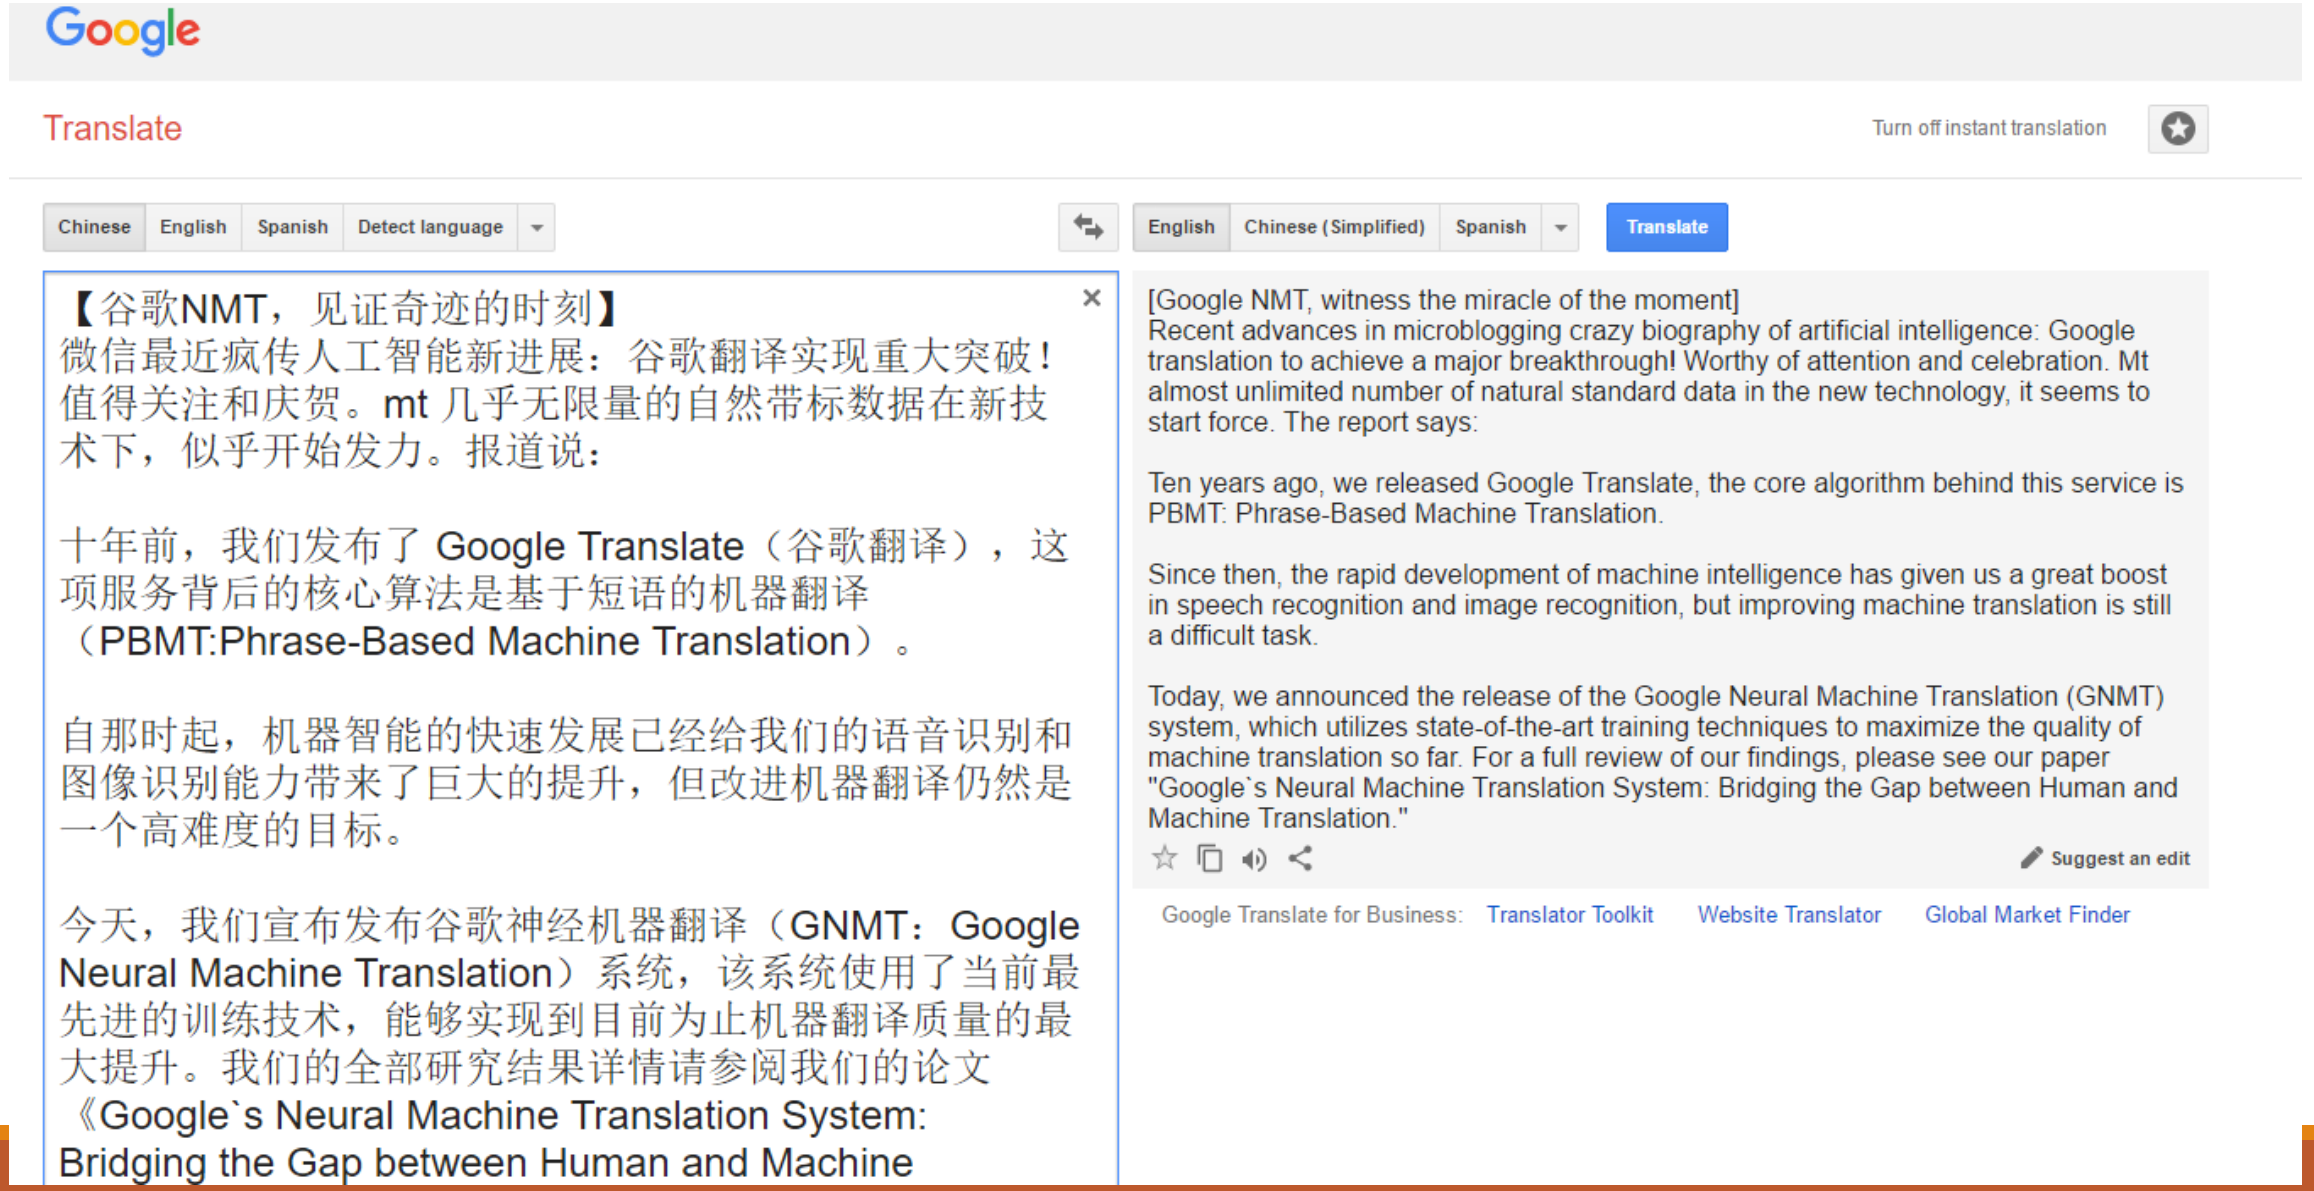

Questioning Answering

- **Goal**: building systems that automatically answer questions posed by humans in a natural language.
- IBM's Watson Jeopardy! (2011), DARPA who/what/where…, Ask Jeeves

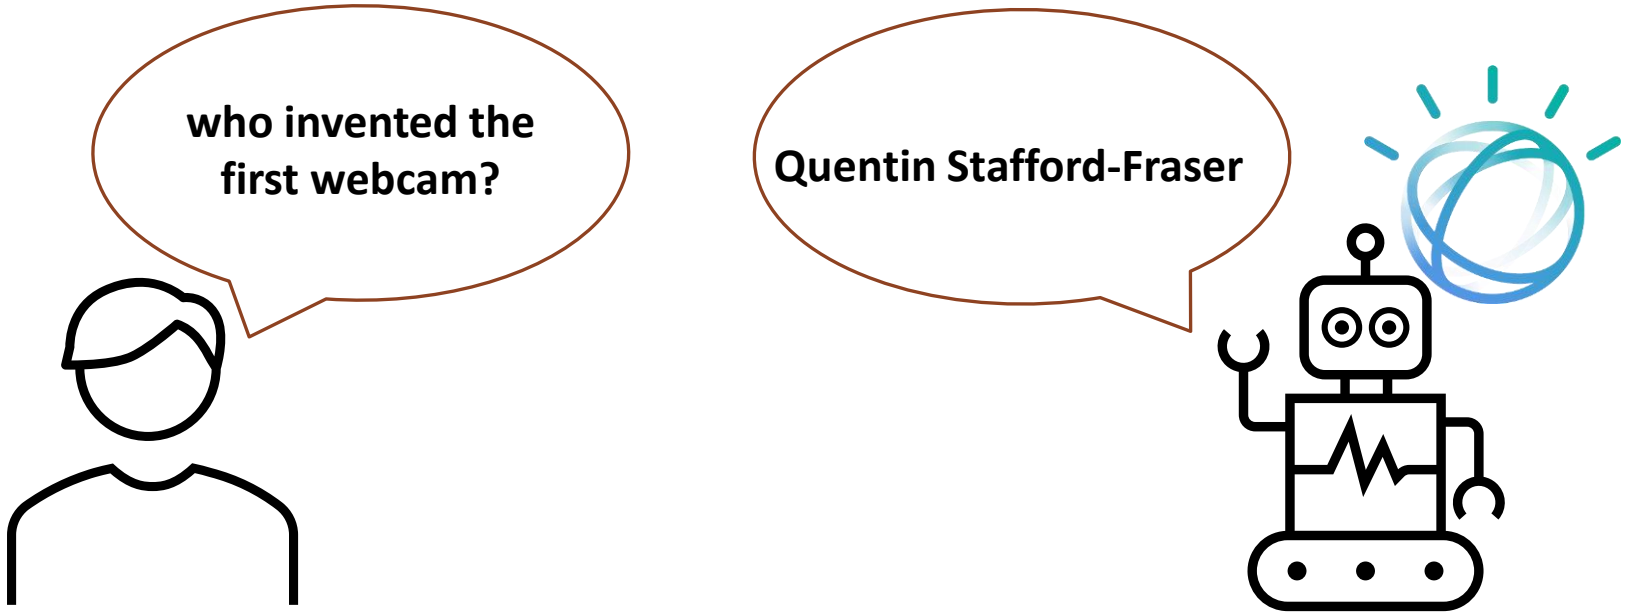

Sentiment Analysis

Discovering people opinions, emotions and feelings about a product or service

http://paxcom.net/blog/why-is-customer-sentiment-analysis-important-for-your-brand/

Sentiment Analysis

- Wow, great place!
- Wow, 35 minutes to get a cup of coffee? Great job.
- Not great but works as expected.
- At first I hated it, but once the story hooked me, I found it difficult to put the book down

Information Extraction

- **Goal**: Extracting specific information from textual sources.

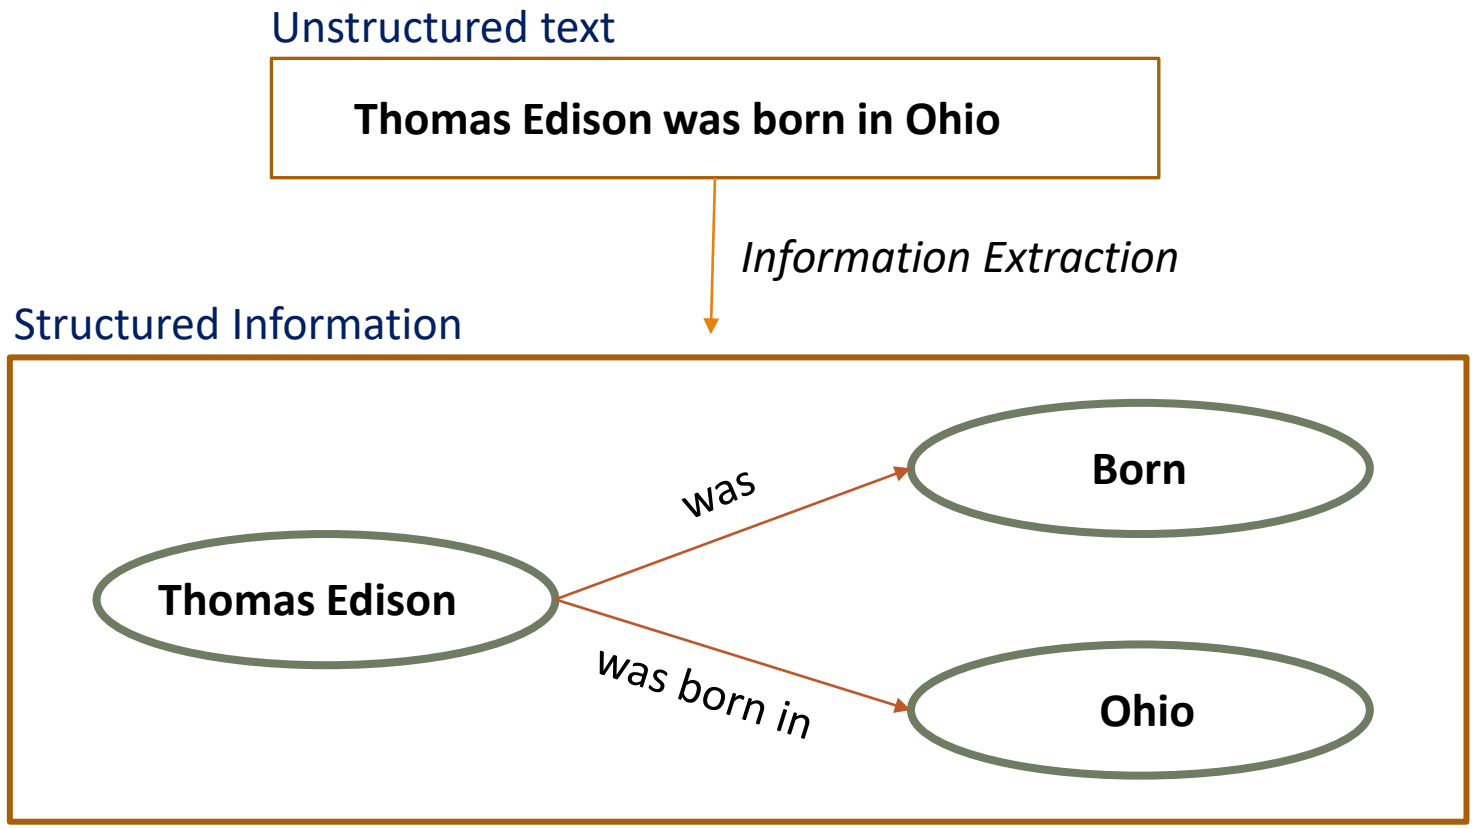

Dialogue Systems

- Goal: A computer system intended to converse with a human.
- Example: Amtrak's 'Julie' and Google Assistant

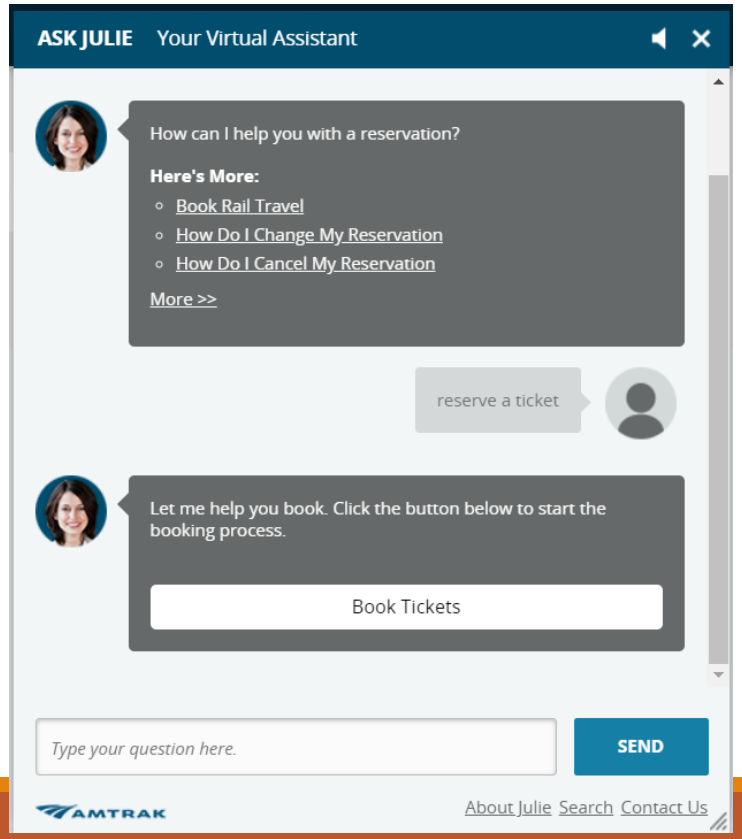

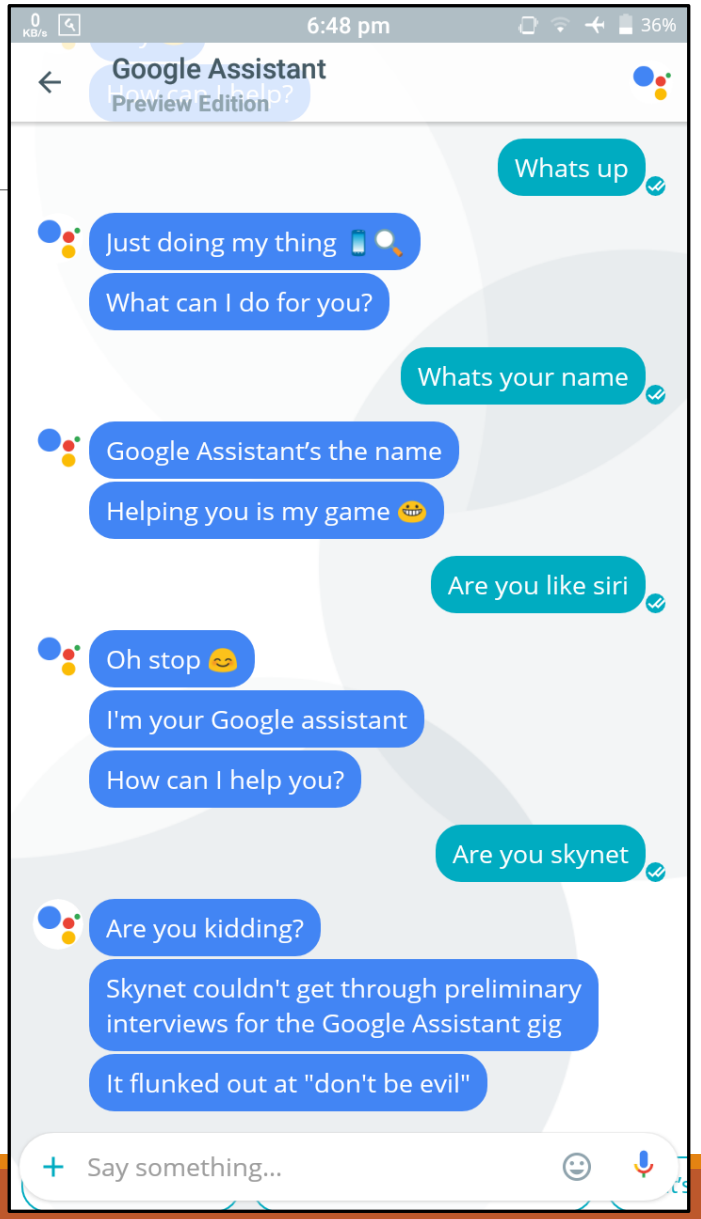

Summarization

- **Goal**: Summarizing very large amounts of text or speech: e.g. your email, the news, voicemail
- Example: https://www.agolo.com

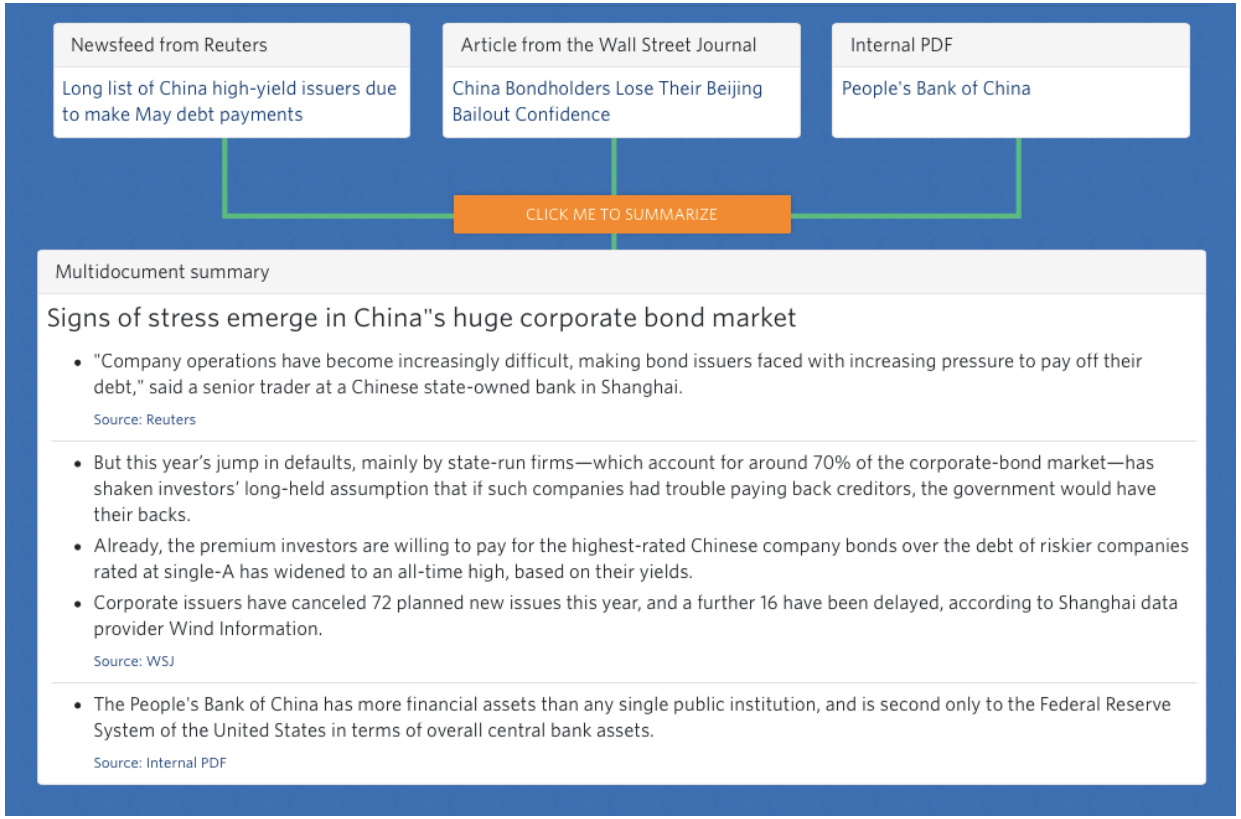

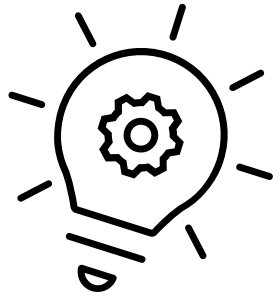

What other NLP applications that you can think of?

Is NLP difficult?

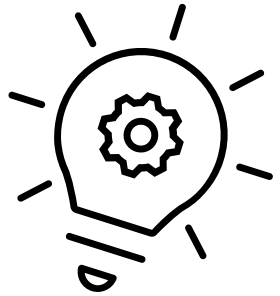

What Makes NLP Difficult?

- **The nature of languages** → Language encode meaning. Language is learned intuitively
- Non-standard text
  - " we're soooo proud of u!"
- Idioms and metaphors
  - "dark horse" "cold feet" "lose face"
- Segmentation
  - "The New York-New Haven railroad"
- Named entities
  - "*Let It Be* sold millions"
- **Ambiguity**

Language Variation

Ambiguity

- Language is ambiguous
- Ambiguity involves multiple or alternative linguistic structures
- Ambiguity results from the existence of multiple possibilities for linguistic levels
- All 6 levels of linguistic knowledge require resolving ambiguity

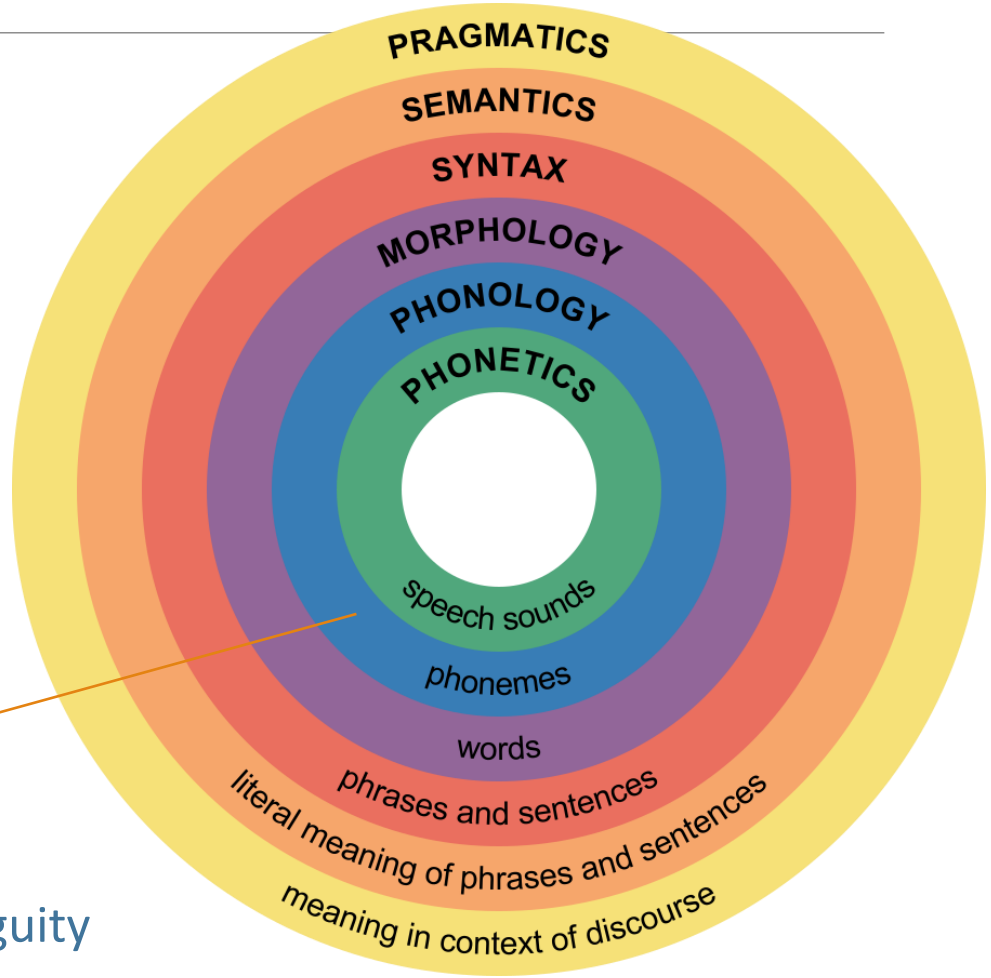

-two, too, to
-ice cream, I scream

➢ Phonological Ambiguity

Ambiguity

Morphological Ambiguity

Unlockable: [[un-lock]-able]
[un-[lock-able]]

Syntactic Ambiguity

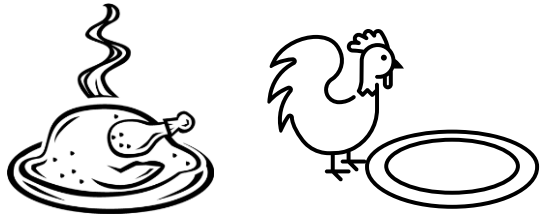

The chicken is ready to eat

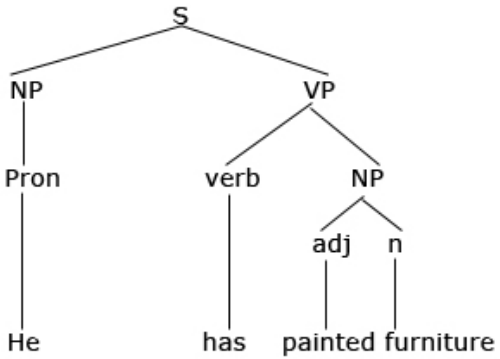

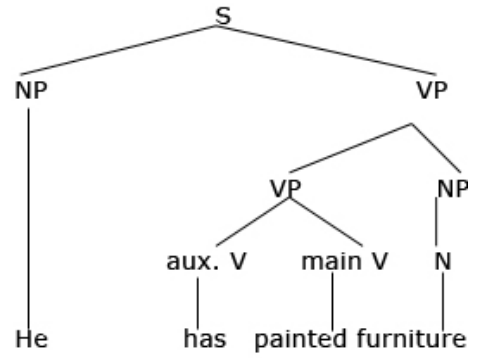

Lexical Ambiguity

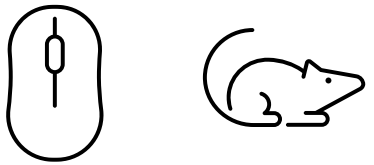

Pass me the mouse

Ambiguity

I made her duck

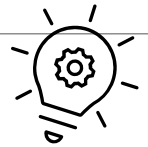

- I cooked waterfowl for her 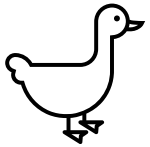
- I cooked the waterfowl that belongs to her
- I created the ceramic duck she owns 
- I caused her to quickly lower her head
- And more….

I made her duck for lunch

I made her duck with clay

I made her duck as the ball was about to hit her

Context really matters

Ambiguity

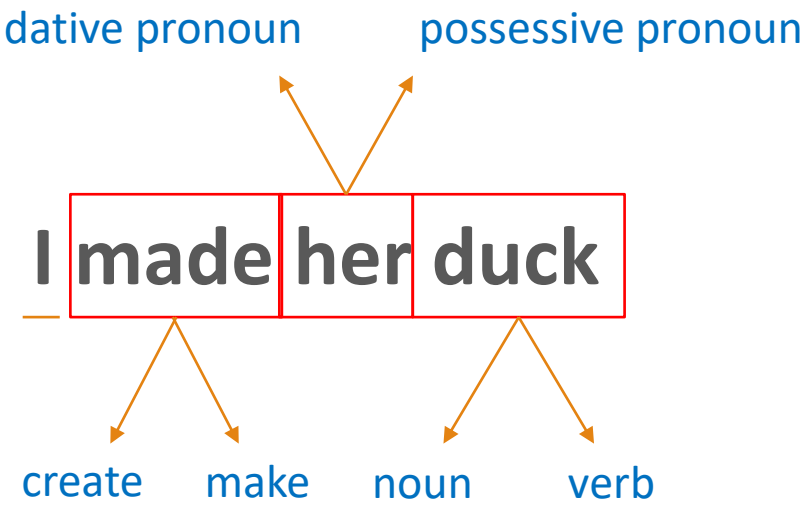
dative pronoun
possessive pronoun
I made her duck
create
make
noun
verb

Dealing with Ambiguity

- Tightly coupled interaction among processing levels; knowledge from other levels can help decide at ambiguous levels.
- Pipeline processing that ignores ambiguity as it occurs and hopes that other levels can eliminate incorrect structures.
- Probabilistic approaches based on making the most likely choices.
- Don’t do anything, maybe it won’t matter.
  - We’ll leave when the duck is ready to eat.
  - The duck is ready to eat now.
    - Does the “duck” ambiguity matter with respect to whether we can leave?

Making Progress …

The task is difficult! What tools do we need?

- Knowledge about language
- Knowledge about the world
- A way to combine knowledge sources

How we generally do this:

- probabilistic models built from language data
  - P(“maison” → “house”) high
  - P(“noir” → “moon”) low

Luckily, rough text features can often do half the job.

Making Progress …

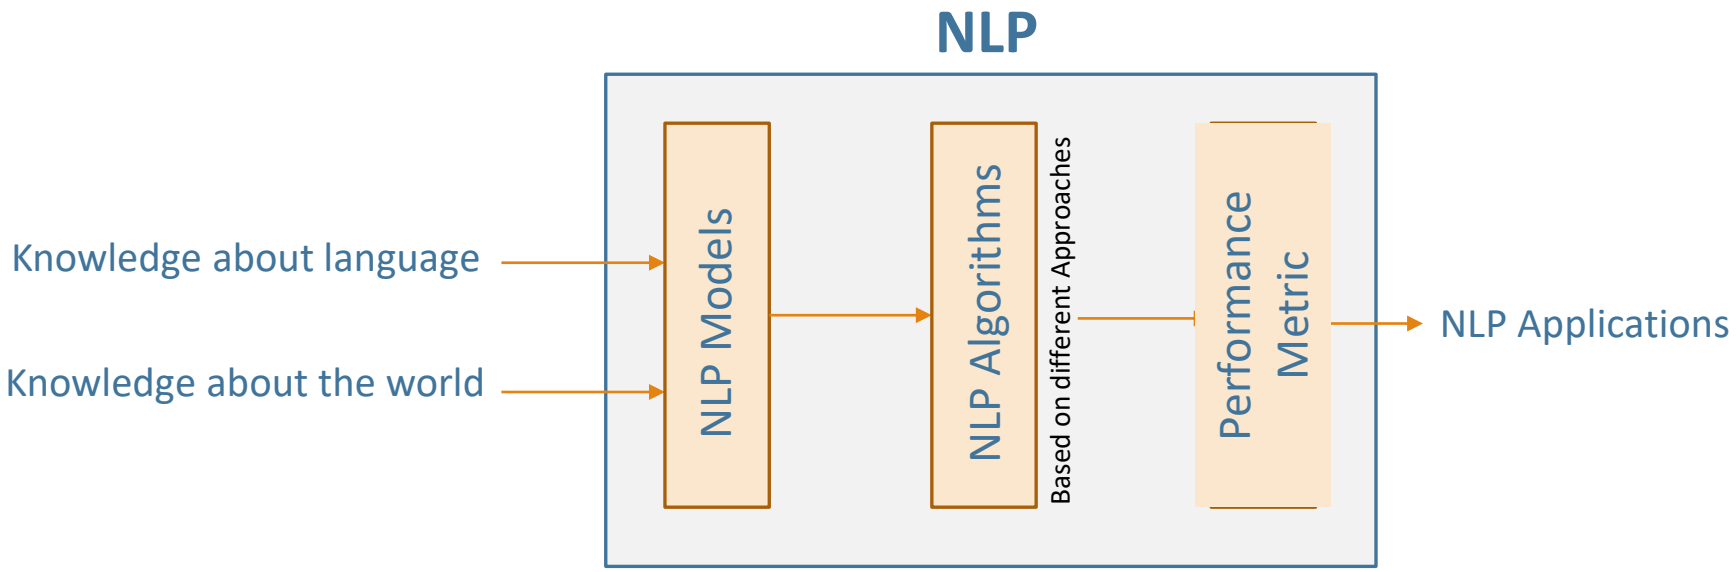

NLP Models

State Machines
- Finite state automata, transducers

Formal Rule Systems
- Regular Grammars, Context Free Grammars

Logic
- First order logic, predicate calculus

Probability Theory
- Associating probabilities with the previous machinery
- Crucial for capturing every kind of linguistic knowledge

Vector-space Models
- An algebraic model for representing text documents (and any objects, in general) as vectors

NLP Algorithms

- State space search algorithms, such as dynamic programming
- Expectation-Maximization (EM)
- Machine learning algorithms, such as classifiers and sequence models, which play a significant role in many language processing tasks

Machine Learning

Machine learning based classifiers that are trained to make decisions based on (implicitly or explicitly modeled) features from context

Simple Classifiers:

- Naïve Bayes
- Logistic Regression
- Decision Trees
- Neural Networks

Sequence Models:

- Hidden Markov Models
- Maximum Entropy Markov Models
- Conditional Random Fields
- Recursive Neural Networks (RNNs, LSTMs)

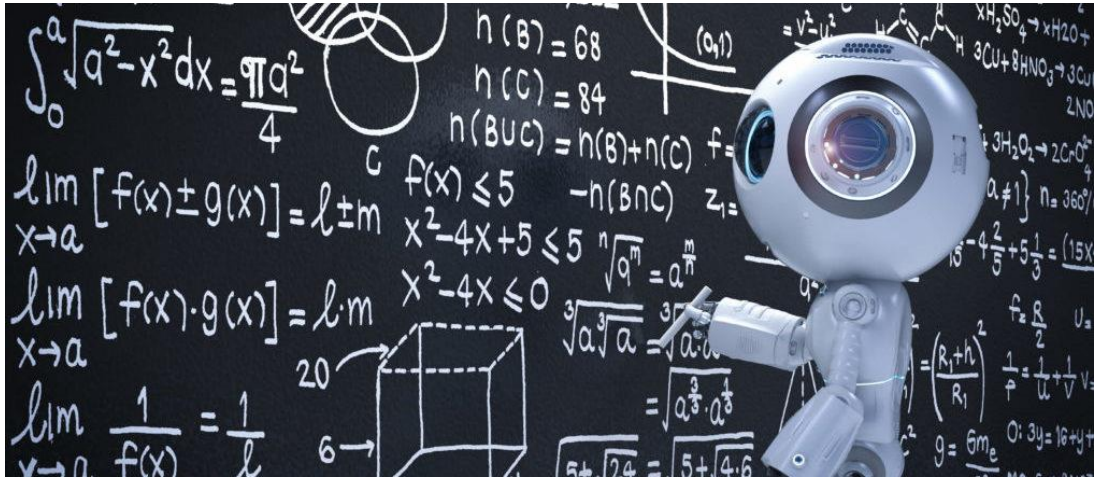

Machines learn from data

NLP Approaches

Rule-based/Symbolic Approaches

- Linguists write rules that are applied by the machines

Corpus-based/Statistical Approaches

- Machines learn the “rules” from training data
  - Annotated data – supervised methods
    - Parallel Corpora: translated text collections
    - Treebanks: manually syntactically analyzed texts
    - Speech Corpora with transcripts
  - Unannotated data – unsupervised methods
  - Semi-supervised methods

Performance Metrics

Methods to evaluate the performance of an NLP system

- Extrinsic Evaluation:
  - Incorporate NLP system into action
- Intrinsic Evaluation:
  - Automatic Evaluation
    - ✓Does system agree with pre-judged examples?
  - Human Post-hoc Evaluation

Historical Notes

1940s - 1950s

- Automata, regular expressions, and Formal Language theory
- Information Theory and foundational research in speech recognition (digits).
- "MT is rather easy" (MIT, Georgetown)

1960s

- "MT is too hard." (ALPAC report)
- Cancelled all work on Machine Translation in the US.
- CL/NLP research starts (e.g. ACL)
- Transformational paradigm in linguistics (Chomsky)

Historical Notes

1970s

- "MT Winter" in US
- Parsing comes of age: CFGs, ATNs,…
- Speech understanding starts

1980s

- Use of probabilistic models in MT (IBM)
- New focus on model evaluation.
- Natural language generation

1990s

- The first search engine using indexing and Google’s success
- More progress

Historical Notes

2000s

- Use of Neural models
- Multitask learning
- Sequence to sequence models
- Attention and pretrained language models

In summary

- Language to Knowledge
  - A lot more to do
- NLP is difficult
  - Ambiguous and various
  - Context really matters
- Machine and deep Learning
  - With enough data and some math computers can do it
  - The future looks exiting for NLP

Outline

- Course Introduction
- **Course Information**
- Deadlines

This Course

Key theory and methods for statistical NLP:

- Finite State Machines and Transducers
- N-gram language modeling
- Hidden Markov Models
- Discriminative classifiers
- Syntactic and Statistical Parsing
- Vector models of meaning
- Word Embedding
- Deep Learning Models

Practical, real-world applications

- Information extraction
- Spelling correction
- Sentiment analysis
- Machine Translation
- Summarization
- Dialogue Systems

Skills you will need

- Basic probability theory
- Simple linear algebra (vectors, matrices)
- Machine Learning Techniques
- Proficiency in Python programming

Logistics

Lectures: Tuesday 06:10PM - 08:40PM

Office Hours: Thursday 11:00AM-1:00PM or by appoitment

Contact**:** raref@gwu.edu

Syllabus: available on BB

Discussion: piazza.com/gwu/fall2020/csci39076907nlp

Text: Mainly from the Jurafsky and Martin 2nd edition and 3rd edition(available online)

Coursework and Grading

- Course Material
  - Course slides and recorded class videos will be available on BB shortly after a class ends
- Assignments
  - Will be posted and announced on Blackboard
  - 4 Programming Assignments and Reading Assignments (Book Chapters and Research Papers)
  - The first assignment will be released next week
  - You have 7 free grace days for the whole semester
- Midterm
  - 10/13
- No Final Exam but a Final Project

Grading

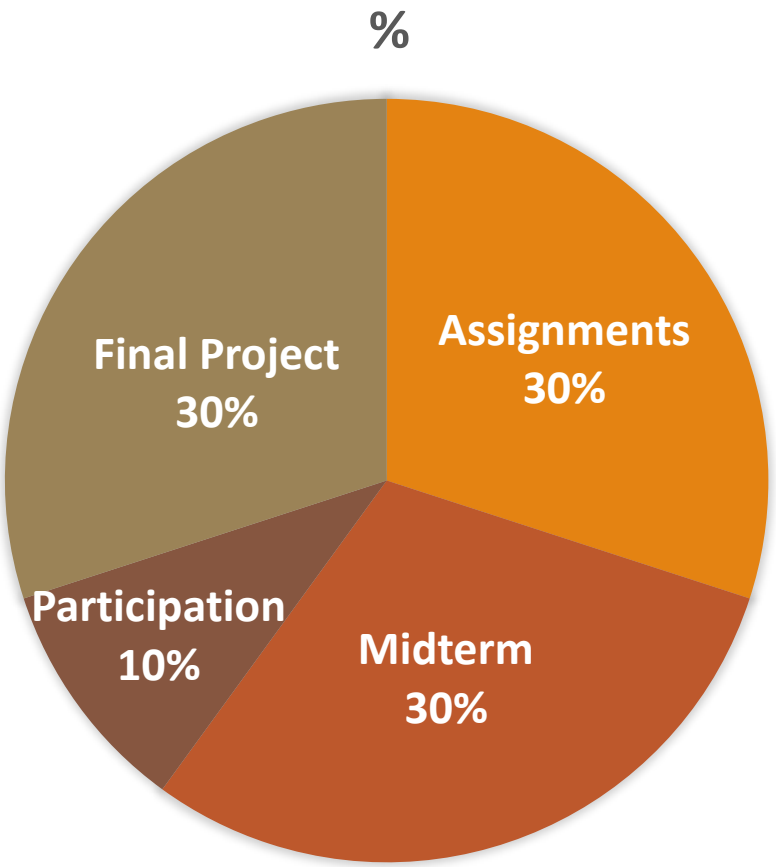
%
Assignments
30%
Midterm
30%
Participation
10%
Final Project
30%

Final Project

- Groups
  - ✓Start forming groups of maximum 4
- Project Topic
  - ✓ Try to think of a novel problem
- Project proposal (20%)
- Class presentation (20%)
- Final report or short paper (30%)
- System Implementation (30%)

Outline

- Course Introduction
- Course Information
- **Deadlines**

Due Dates

9

- Assignment 1, 9/25
- Topic of interest, 9/21
- Project groups 9/29

10

- Assignment 2, 10/9
- Midterm, 10/13
- Project Proposal, 10/30

11

- Assignment 3, 11/13
- Assignment 4, 11/27

12

- Final Project Presentations and Submissions, 12/1 and 12/8

Questions??

Contributions to the course material & slides

Slides are sometimes adapted (with permission) from other great slide sets, namely from:

- Mona Diab, Chris Manning, Dan Jurafsky, Jason Eisner, Rada Mihalcea, Michael Collins, Alessandro Moschitti, Julia Hirschberg, Kathleen McKeown, Dragomir Radev.